



Advanced Operating Systems: Three Easy Pieces

Introduction to Operating Systems

Course Content

- Cover the following advanced topics including:
 - ❑ Virtualization
 - ❑ Memory Management
 - ❑ Concurrency Control and Deadlock
 - ❑ Persistence
 - ❑ Security
 - ❑ Distributed Operating System
- Cover briefly few advanced case studies
- Cover briefly few advanced research topics in Operating Systems
- In addition, students (divided in groups) will work on a set of hands-on projects.

Textbook

- **“Operating Systems: Three Easy Pieces,”** by Remzi H. Arpaci-Dusseau and Andrea C. Arpaci-Dusseau.

Arpaci –Dusseau Books

March, 2015 (version 0.90)

<http://pages.cs.wisc.edu/~remzi/OSTEP/>

Three Pieces are: Virtualization, Concurrency, and Persistence



What Does Operating System Provide?

Operating System Roles?

- **Role #1: Abstraction to applications that is hardware independent of the physical hardware resources**
- **What is a resource and its abstraction?**

Physical Resource	Abstraction
CPU	Process / Thread
Memory	Address Space
Disk	Files



Operating System Roles?

■ Role #2: Resource Management

1. Protect applications from each other
2. Provide efficient access to resources
3. Provide fair access to resources



Operating System Organization

OS Organization: Three Pieces?

■ First Piece: Virtualization

- ❑ Make each application believe it has each resources to itself

■ Second Piece: Concurrency

- ❑ Correct behavior while events are occurring simultaneously and may interact with one another

■ Third Piece: Persistence

- ❑ Access information permanently – lifetime of data is longer and independent of lifetime of any one process / application even in the presence of failure/reboot

OS Organization: Three Pieces?

■ Advanced Topics:

- ❑ Multiprocessor (tightly coupled)
- ❑ Distributed Systems (loosely coupled)
- ❑ Virtual Machine
- ❑ NUMA and RDMA
- ❑ Data Center OS (DC OS)



First Piece: Virtualization

What happens when a program runs?

■ A running program executes instructions:

1. The processor **fetches** an instruction from memory.
2. **Decode**: Figure out which instruction this is
3. **Execute**: i.e., add two numbers, access memory, check a condition, jump to function, and so forth.
4. The processor moves on to the **next instruction** and so on.

Operating System (OS)

■ Responsible for:

- ❑ Making it easy to **run** programs
- ❑ Allowing programs to **share** memory
- ❑ Enabling programs to **interact** with devices

OS is in charge of making sure the system operates **correctly** and **efficiently**.

Virtualization

- The OS takes a **physical resource** and transforms it into a **virtual form** of itself.
 - ❑ **Physical resource:** Processor, Memory, Disk ...
 - ❑ **The virtual form** is more general, powerful and easy-to-use.
 - ❑ Sometimes, we refer to the OS as a **virtual machine**.

System call

- **System call** allows user to ask the OS for what to do or what services is needed:
 - The OS provides some interface (APIs, standard library) to provide **system-level services**.
 - A typical OS exports a few hundreds system calls.
 - Run programs
 - Access memory
 - Access devices

The OS is a resource manager.

- The OS **manages resources** such as *CPU*, *memory* and *disk*.
- **The OS allows:**
 - ❑ Many programs to run → Sharing the CPU
 - ❑ Many programs to *concurrently* access their own instructions and data → Sharing memory
 - ❑ Many programs to access devices → Sharing disks

Virtualizing the CPU

- The system has a very large number of virtual CPUs.
 - Turning a single CPU into a seemingly infinite number of CPUs.
 - Allowing many programs to seemingly run at once
→ **Virtualizing the CPU**

Virtualizing the CPU (Cont.)

```
1      #include <stdio.h>
2      #include <stdlib.h>
3      #include <sys/time.h>
4      #include <assert.h>
5      #include "common.h"
6
7      int
8      main(int argc, char *argv[])
9      {
10         if (argc != 2) {
11             fprintf(stderr, "usage: cpu <string>\n");
12             exit(1);
13         }
14         char *str = argv[1]; // str = 1st cmd line arg
15         while (1) {
16             Spin(1); // Repeatedly checks the time and
17                     // returns once it has run for a second
18             printf("%s\n", str);
19         }
20         return 0;
21     }
```

Simple Example(cpu.c): Code That Loops and Prints

Virtualizing the CPU (Cont.)

■ Execution result 1:

```
prompt> gcc -o cpu cpu.c -Wall
prompt> ./cpu "A"
A
A
A
^C
prompt>
```

Run forever; Only by pressing "Control-c" can we halt/exit the program

Virtualizing the CPU (Cont.)

■ Execution result 2:

```
prompt> ./cpu A & ; ./cpu B & ; ./cpu C & ; ./cpu D &  
[1] 7353  
[2] 7354  
[3] 7355  
[4] 7356  
A  
B  
D  
C  
A  
B  
D  
C  
A  
C  
B  
D  
...
```

Even though we have only **one processor**, all four of programs seem to be running **at the same time!**

Virtualizing Memory

- The physical memory is an array of bytes.
- A program keeps all of its data structures in memory:
 - **Read memory** (load):
 - Specify an address to be able to access the data
 - **Write memory** (store):
 - Specify the data to be written to the given address

Virtualizing Memory (Cont.)

```
1      #include <unistd.h>
2      #include <stdio.h>
3      #include <stdlib.h>
4      #include "common.h"
5
6      int
7      main(int argc, char *argv[])
8      {
9          int *p = malloc(sizeof(int)); // a1: allocate some memory
                                         // to be used as counter
10
11          assert(p != NULL);
12          printf("(%d) address of p: %08x\n",
13                getpid(), (unsigned) p); // a2: print out the pid
                                         // and address of the memory
14          *p = 0; // a3: put zero into the first slot of the allocated memory
15          while (1) {
16              Spin(1);
17              *p = *p + 1;
18              printf("(%d) p: %d\n", getpid(), *p); // a4: print
                                                         pid and the content of memory counter
19          }
20          return 0;
21      }
```

Virtualizing Memory (Cont.)

■ The output of the program mem.c

```
prompt> ./mem
(2134) memory address of p: 00200000
(2134) p: 1
(2134) p: 2
(2134) p: 3
(2134) p: 4
(2134) p: 5
^C
```

- ❑ The newly allocated memory is at address 00200000.
- ❑ It updates the value and prints out the result.

Virtualizing Memory (Cont.)

■ Running mem.c multiple times:

```
prompt> ./mem & ./mem &  
[1] 24113  
[2] 24114  
(24113) memory address of p: 00200000 // Heap address  
(24114) memory address of p: 00200000 // Heap address  
(24113) p: 1  
(24114) p: 1  
(24114) p: 2  
(24113) p: 2  
(24113) p: 3  
(24114) p: 3  
...
```

- It is as if each running program has its **own private memory**.
 - Each running program has allocated memory at the same address.
 - Each seems to be updating the value at 00200000 independently.

Virtualizing Memory (Cont.)

- Each process **accesses** its own private **virtual address space**:
 - ❑ The OS maps **address space** onto the **physical memory**.
 - ❑ A memory reference within one running program does not affect the address space of other processes.
 - ❑ Physical memory is a shared resource, managed by the OS.



Second Piece: Concurrency



The problem of Concurrency

- The OS is juggling **many things at once**, first running one process, then another, and so forth.
- Modern **multi-threaded programs** also exhibit the concurrency problem.

Concurrency Example:

A Multi-threaded Program (thread.c)

```
1      #include <stdio.h>
2      #include <stdlib.h>
3      #include "common.h"
4
5      volatile int counter = 0;
6      int loops;                                /* user supplied value */
7
8      void *worker(void *arg) {
9          int i;
10         for (i = 0; i < loops; i++) {
11             counter++;
12         }
13         return NULL;
14     }
15
16     int
17     main(int argc, char *argv[])
18     {
19         if (argc != 2) {
20             fprintf(stderr, "usage: threads <value>\n");
21             exit(1);
22         }
```

Concurrency Example (Cont.)

```
23         loops = atoi(argv[1]);
24         pthread_t  p1, p2;
25         printf("Initial value : %d\n", counter);  // counter = 0
26
27  /* pthread_create(thread_t t, pthread_attr_t *attr,
28                  void *(*start_routine)(void *), void *arg); */
29         Pthread_create(&p1, NULL, worker, NULL);
30         Pthread_create(&p2, NULL, worker, NULL);
31         Pthread_join(p1, NULL);
32         Pthread_join(p2, NULL);
33         printf("Final value : %d\n", counter);
34         return 0;
35     }
```

□ The main program creates **two threads**:

- **Thread**: a function running within the same memory space. Each thread start running in a routine called worker().
- **worker()**: increments a counter

Concurrency Example (Cont.)

- **loops** determines how many times each of the two workers will **increment the shared counter** in a loop.

❑ loops: 1,000.

```
prompt> gcc -o thread thread.c -Wall -pthread
prompt> ./thread 1000
Initial value : 0
Final value : 2000
```

❑ loops: 100,000.

```
prompt> ./thread 100000
Initial value : 0
Final value : 143012          // huh?? Not 200,000
prompt> ./thread 100000
Initial value : 0
Final value : 137298          // what is this - different value?
```

- ❑ The issue is “loops” is global variable and it is being accessed without synchronization by 2-threads (read-increment-write) is not atomic.

Why is this happening?

- **Increment a shared counter** → take three instructions:
 1. Load the value of the counter from memory into register.
 2. Increment it
 3. Store it back into memory
- These three instructions do not execute **atomically** → **Problem of concurrency** happen.



Third Piece: Persistence

Persistence

- Devices such as DRAM store values/data in a volatile medium.
- *Hardware* and *software* are needed to store data persistently.
 - **Hardware:** I/O device such as a hard drive, solid-state drives (SSDs)
 - **Software:**
 - File system manages the disk.
 - File system is responsible for storing any files the user creates.

Persistence (Cont.)

- Create a file (/tmp/file) that contains the string “hello world”

```
1      #include <stdio.h>
2      #include <unistd.h>
3      #include <assert.h>
4      #include <fcntl.h>
5      #include <sys/types.h>
6
7      int
8      main(int argc, char *argv[])
9      {
10         int fd = open("/tmp/file", O_WRONLY | O_CREAT
                        | O_TRUNC, S_IRWXU);
11         assert(fd > -1);
12         int rc = write(fd, "hello world\n", 13);
13         assert(rc == 13);
14         close(fd);
15         return 0;
16     }
```

open(), write(), and close() system calls are routed to the part of OS called the file system, which handles the requests

Persistence (Cont.)

- **What OS does in order to write to disk?**
 - ❑ Figure out **where** on disk this new data will reside
 - ❑ **Issue I/O** requests to the underlying storage device
- **File system** handles system crashes during write:
 - ❑ **Journaling** or **copy-on-write**
 - ❑ Carefully ordering writes to disk



Operating System Design Goals

OS Design Goals

- Build up **abstraction**:
 - Make the system convenient and easy to use.
 - User program to hardware independent APIs.
- Provide high **performance**:
 - Minimize the overhead of the OS.
 - OS must strive to provide virtualization without excessive overhead.
- **Protection** between applications:
 - **Isolation**: Bad behavior of one does not harm other and the OS itself.

Design Goals (Cont.)

- **High degree of reliability**

- ☐ The OS must also run non-stop.

- **Other issues**

- ☐ Energy-efficiency
- ☐ Security
- ☐ Mobility



END